

ZYVORAI LABS · CUSTOMER DOCUMENTATION

Zyvor Janus

Getting Started

2026-08-29

Zyvor Janus replays GPU schedulers against a cluster config — no physical GPUs required. This page gets you to a first CLI run and an optional web console session.

Prerequisites

Requirement	Notes
Rust toolchain	<code>cargo</code> to build <code>zyvor-janus-cli</code>
Cluster config YAML	<code>configs/clusters/</code> or a Forge CRD export bundle
Node.js (optional)	Only for the Next.js web dashboard

1. Clone and build

```
git clone https://github.com/zyvoral/janus.git
```

```
cd zyvor-janus # or your checkout path
```

```
cargo build -p zyvor-janus-cli
```

2. First simulation (CLI)

```
cargo run -p zyver-janus-cli -- run --config configs/clusters/small-h100.yaml
```

Scheduler choice (FIFO, priority, preemptive, best-fit) comes from the config. LLM-serving workloads also report TTFT/TPS.

3. Optional — start the web console

```
./scripts/run_web_dashboard.sh
```

URL	Service
<code>http://<host>:3000</code>	Next.js UI
<code>http://<host>:8080/api/health</code>	FastAPI backend

Sign in (**Admin / Admin@321** by default) → [Dashboard](#) → **Run**.

Override ports: `API_PORT=9000 UI_PORT=3001 ./scripts/run_web_dashboard.sh`.

4. Orient yourself (UX)

1. Sidebar: Dashboard · Runs · Benchmark · What-if · Twins.
2. Launch a run from Dashboard; open **Run detail** then **Simulate** to replay decisions.
3. Use **What-if** to sweep schedulers; **Benchmark** for TTFT/TPS score vectors.

Troubleshooting

- `cargo run` can't find the config — `--config` is relative to your cwd under `configs/clusters/` .
- **No TTFT/TPS** — only LLM-serving workloads emit those metrics.
- **Dashboard empty** — confirm FastAPI health and that CLI/UI share the same outputs path.

Next steps

- [Using the Dashboard](#)
 - [Configuration](#)
 - [Page guides](#)
 - [Workflows](#)
-

Zyvor Janus ships a CLI and a Next.js web console. Both consume the same simulation engine output.

Shell chrome

Element	Job
Sidebar / glass icons	Dashboard · Runs · Benchmark · What-if · Twins
Top bar brand	Jump home
Collapse control	Icon-only sidebar
Mobile bottom nav	Same primary destinations

Surfaces

Need	Where
Launch / compare configs	Dashboard
Live scheduling replay	Simulate
TTFT/TPS benchmarks	Benchmark
Scheduler sweeps	What-if
AIPerf twins	Twins
One run	Run detail

Operate tips

1. Start the stack with `./scripts/run_web_dashboard.sh` , then sign in.
2. Prefer Dashboard **Run** for interactive work; use CLI for CI and scripts.
3. After a run, **Open in Simulate** before changing schedulers in What-if.
4. Never paste lab IPs into runbooks — use `<host>` .

Related

- [Getting Started](#)
 - [Page index](#)
 - [Admin basics](#)
-

First CLI run

1. [Getting Started](#) — build + `run --config configs/clusters/small_h100.yaml`
2. Confirm the run completes; note TTFT/TPS only for LLM-serving configs

Console: run → simulate → inspect

1. `./scripts/run_web_dashboard.sh` → [Login](#)
2. [Dashboard](#) → pick config → **Run**
3. Open [Run detail](#) → **Open in Simulate**
4. Watch Accept→assign decisions on [Simulate](#)

Compare two configs

1. Dashboard → **Compare configs** (A vs B)
2. Or CLI two runs + Benchmark / What-if for scheduler sweeps

Scheduler what-if sweep

1. [What-if](#) → select schedulers → **Run sweep**
2. Download CSV; open best candidate run detail

Benchmark with twins

1. Calibrate / list [Twins](#)
2. [Benchmark](#) → **Run benchmark** → read sim vs measured

Related

- [Configuration](#)
- [Using the Dashboard](#)