

ZYVORAI LABS · CUSTOMER DOCUMENTATION

Zyvor Janus

Page-by-Page Product Manual

2026-08-29

Page-by-page guides

Each guide follows: Purpose → When to use it → How to get there → Operate from the console (UX) → Related pages.

Every route is also listed in the [complete page index](#).

Auth

Page	What it covers
Login	Sign in to the Zyvor Janus web console (session cookie).

Console

Page	What it covers
Benchmark	Scheduler benchmarks with inference TTFT/TPS metrics, score vectors, and AIPerf twin comparison.
Dashboard	Mission Control — launch simulations, compare configs, and jump into recent runs.
Run detail	Inspect one simulation run — cluster/queue view, metrics, and jump to Simulate.
Simulate	Live scheduling replay — watch requests → queue → GPU assign → complete for a run.
Twins	AIPerf-calibrated GPU/model performance twins that ground the simulator against measured hardware.
What-if	Sweep schedulers on the same inference workload and compare a results matrix.

7 guides. Regenerate: `node scripts/customer-docs/generate-guide-index.mjs` .

Login

Purpose

Sign in to the Zyvor Janus web console (session cookie).

When to use it

- Operate **Login** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/login`
- Nav: **Unauthenticated** → `/login`

Operate from the console (UX)

1. Open `http://<host>:3000/login` (or you are redirected here).
2. Enter dashboard credentials (default **Admin / Admin@321** unless `ZYVOR_JANUS_DASHBOARD_USER` / `PASSWORD` override).
3. Submit — session is an `httpOnly` cookie (≈ 7 days).
4. **Empty / fail:** Login failed → check API on `:8080` and env credentials; middleware blocks `/api` without auth.
5. **Success:** Land on Dashboard; sidebar shows Dashboard · Runs · Benchmark · What-if · Twins.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Dashboard](#)
 - [Getting Started](#)
 - [Page index](#)
-

Benchmark

Purpose

Scheduler benchmarks with inference TTFT/TPS metrics, score vectors, and AIPerf twin comparison.

When to use it

- Operate **Benchmark** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/benchmark`
- Nav: **Sidebar** → **Benchmark**

Operate from the console (UX)

1. Open `/benchmark`.
2. Configure the benchmark case → **Run benchmark**.
3. Read Score vector radar and Sim vs measured (AIPerf twin) panels.
4. Open recent benchmark reports → link into `/runs/:id`.
5. **Empty / fail:** Running stuck → API/engine error; no TTFT → non-LLM workload.
6. **Success:** Report row appears with colored cells vs best-in-column.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Run detail](#)
 - [Twins](#)
 - [Dashboard](#)
 - [Getting Started](#)
 - [Page index](#)
-

Dashboard

Purpose

Mission Control — launch simulations, compare configs, and jump into recent runs.

When to use it

- Operate **Dashboard** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/`
- Nav: **Sidebar** → **Dashboard**

Operate from the console (UX)

1. Open `/` after sign-in.
2. **Launch simulation**: pick a cluster config → optional shadow scheduler → **Run** (or **Run + shadow**).
3. **Compare configs**: choose A and B → Compare → review metric deltas in ComparePanel.
4. Open a recent run → `/runs/:id`, or jump to Benchmark / What-if shortcuts.
5. Anchor `/#runs` focuses the runs list (sidebar Runs).
6. **Empty / fail**: No configs → ensure `configs/clusters/` is visible to the API; Run fails → check FastAPI logs on `:8080`.
7. **Success**: New run appears; Open navigates to run detail.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Run detail](#)
 - [Benchmark](#)
 - [What-if](#)
 - [Simulate](#)
 - [Getting Started](#)
 - [Page index](#)
-

Run detail

Purpose

Inspect one simulation run — cluster/queue view, metrics, and jump to Simulate.

When to use it

- Operate **Run detail** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/runs/:id`
- Nav: **Dashboard runs** · **Benchmark reports** → **Open**

Operate from the console (UX)

1. Open `/runs/<id>` from a run link.
2. Read status, metrics, and cluster/queue panels.
3. **Open in Simulate** to replay scheduling decisions.
4. **Back to dashboard** when done.
5. **Empty / fail:** Run not found → stale id or API offline.
6. **Success:** Run payload loads; Simulate deep-link works.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Simulate](#)
 - [Dashboard](#)
 - [Benchmark](#)
 - [Getting Started](#)
 - [Page index](#)
-

Simulate

Purpose

Live scheduling replay — watch requests → queue → GPU assign → complete for a run.

When to use it

- Operate **Simulate** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/simulate`
- Nav: **Run detail** → **Open in Simulate** · `/simulate?run=<id>`

Operate from the console (UX)

1. Open `/simulate` or `/simulate?run=<id>` from a completed/in-progress run.
2. Select a run if needed → **Watch** / load to stream the SimulateStage.
3. Follow legend for scheduler decisions; respect reduced-motion preferences.
4. Jump **Run detail** or **Dashboard** from actions.
5. **Empty / fail**: No runs to watch → start one from Dashboard; busy spinner → wait for snapshot.
6. **Success**: Stage animation advances with decisions for the selected run.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Run detail](#)
 - [Dashboard](#)
 - [Getting Started](#)
 - [Page index](#)
-

Twins

Purpose

AIPerf-calibrated GPU/model performance twins that ground the simulator against measured hardware.

When to use it

- Operate **Twins** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/twins`
- Nav: **Sidebar** → **Twins**

Operate from the console (UX)

1. Open `/twins`.
2. Browse calibrated twins; sort/filter by `measured_at`.
3. Open linked AIPerf run when `aiperf_run_id` is present.
4. **Empty / fail:** No twins calibrated yet → run AIPerf calibration pipeline per product docs.
5. **Success:** Twin cards list hardware/model pairs used by Benchmark sim-vs-measured.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Benchmark](#)
 - [Run detail](#)
 - [Getting Started](#)
 - [Page index](#)
-

What-if

Purpose

Sweep schedulers on the same inference workload and compare a results matrix.

When to use it

- Operate **What-if** when your job matches this screen
- Prefer **Dashboard** after login if you are unsure where to start
- Confirm FastAPI health on `:8080` if the UI looks empty

How to get there

- Route: `/what-if`
- Nav: **Sidebar** → **What-if**

Operate from the console (UX)

1. Open `/what-if`.
2. Pick base config (e.g. `small_h100.yaml`) and subset of schedulers.
3. **Run sweep** → inspect Results matrix bars; **Download CSV** when offered.
4. **Empty / fail**: No schedulers selected → enable at least one; sweep fails → config path.
5. **Success**: Matrix populated for each scheduler under test.

Web UI: `http://<host>:3000` · API: `http://<host>:8080/api/health` (`./scripts/run_web_dashboard.sh`).

CLI: `cargo run -p zyvor-janus-cli -- run --config ...` . Never publish lab IPs — use `<host>` .

Related pages

- [Dashboard](#)
 - [Benchmark](#)
 - [Getting Started](#)
 - [Page index](#)
-

Zyvor Janus — Complete page index

Every primary navigable dashboard route.

Generated: 2026-08-29 · 7 routes

Regenerate: `node scripts/customer-docs/generate-page-index.mjs`

AUTH

Page	Route	Purpose	Guide
Login	<code>/login</code>	Sign in to the Zyvor Janus web console (session cookie).	Open

CONSOLE

Page	Route	Purpose	Guide
Dashboard	<code>/</code>	Mission Control — launch simulations, compare configs, and jump into recent runs.	Open
Simulate	<code>/simulate</code>	Live scheduling replay — watch requests → queue → GPU assign → complete for a run.	Open
Benchmark	<code>/benchmark</code>	Scheduler benchmarks with inference TTFT/TPS metrics, score vectors, and AIPerf twin comparison.	Open
What-if	<code>/what-if</code>	Sweep schedulers on the same inference workload and compare a results matrix.	Open
Twins	<code>/twins</code>	AIPerf-calibrated GPU/model performance twins that ground the simulator against measured hardware.	Open
Run detail	<code>/runs/:id</code>	Inspect one simulation run — cluster/queue view, metrics, and jump to Simulate.	Open

Related

- [Customer docs home](#)
- [Page-by-page guides](#)